

Spectral Energy is Not a Reliable Collapse Diagnostic for Low-Resource Quantized LoRA: Evidence from Turkic Languages

Anonymous ARR Submission

Abstract

We study the SVD dynamics of LoRA adapters fine-tuned on three low-resource Turkic languages (Kyrgyz, Kazakh, Uzbek) using Gemma-2-9B (4-bit NF4 and BF16), and report three findings.

(1) The spectral-energy threshold of Biderman et al. (2024) silently misses learning-rate-induced collapse. A controlled three-way decomposition (E5/E5b/E5c) shows three $r=64$ runs converging to near-identical settled SE trajectories (final max SE 0.06–0.12 in all three) but spanning the full functional range from catastrophic collapse (NER F1 = 0, TUMLU $\sim 23\%$, cross-lingual PPL > 600) to baseline-better target performance (KY PPL 3.86 vs. 4.78 for $r=16$). The collapse is purely learning-rate-induced; SE never approaches the proposed 0.7 threshold. A BF16 control (C3, $n=1$) rules out 4-bit quantization as the cause: BF16 SE is statistically indistinguishable from 4-bit at every comparable point.

(2) Frobenius-norm growth is a useful but language-dependent diagnostic. It separates regimes *within* an underrepresented-language category ($1.13\times$ for transfer warm-start vs. $15\text{--}38\times$ for pathological runs) but fails across categories: EN grows $33.3 \pm 0.1\times$ across two seeds without forgetting (KZ PPL 5.52 ± 0.02), comparable to Kazakh’s $38\times$ with severe forgetting. We therefore recommend tracking $\|B\|_F$ trajectories together with cross-lingual PPL probes, not fixed thresholds.

(3) Related-language warm-start preserves source-language knowledge. Sequential Kazakh $\rightarrow$ Kyrgyz initialization yields a $2\times$ reduction in cross-lingual Kazakh PPL (23.5 ± 0.1 vs. 50.0 ± 2.4 across two seeds, $\sim 10\times$ outside seed variance),

with only $+13\%$ Frobenius-norm growth. A cross-family control (EN $\rightarrow$ KY) rules out a generic warm-start explanation: starting from an unrelated language yields KZ PPL 47.46, essentially identical to direct training (47.58). The effect is specific to related languages. The single-seed NER-F1 advantage ($+61\%$) does not survive reseeding; we frame the effect as KZ retention rather than KY improvement.

Scope. Seven main configurations are $n=2$ seed-replicated (E1b, E3, E5, E5b, E5c, E6, E8); five (E1, E2, E4, C3, E7) are $n=1$. Numbers are reported as consistent-across-two-seeds rather than population-variance estimates. All code, adapters, and the BF16 control bundle are released.

1 Introduction

Low-Rank Adaptation (Hu et al., 2022) has become the dominant approach for adapting large language models (LLMs) to new domains and languages, particularly when computational resources are limited. By decomposing weight updates into low-rank matrices $\Delta W = BA$, LoRA achieves parameter efficiency while preserving most of the pretrained model’s capabilities. However, the internal dynamics of LoRA weight matrices during training—specifically their spectral properties—remain poorly understood, especially for morphologically complex, low-resource languages.

Recent work has identified *spectral collapse* as a potential failure mode of LoRA fine-tuning, where adapter weight matrices degenerate into rank-one approximations dominated by a single singular value (Biderman et al., 2024; Jang et al., 2024). Biderman et al. (2024) proposed that when the spectral energy (SE)—the ratio of the largest singular value to the sum of all singular values—exceeds a threshold of approximately 0.7, the adapter

has effectively collapsed and further training becomes unproductive. However, these findings were established primarily on English-language tasks with relatively large training corpora.

Turkic languages present a particularly challenging testbed for LoRA adaptation. Kyrgyz, Kazakh, and Uzbek are agglutinative languages with rich morphological systems, employ both Cyrillic and Latin scripts, and suffer from severe underrepresentation in multilingual LLM pretraining data. Critically, they exhibit extreme subword fragmentation under standard tokenizers ($3.2\times$ more tokens per word than English), which forces LoRA adapters to make larger weight modifications to accommodate poorly-covered morphology. These properties make Turkic language adaptation a stress test for parameter-efficient fine-tuning methods.

In this work, we systematically investigate the spectral dynamics of LoRA adapters trained on three Turkic languages and test two hypotheses:

- **H1 (Spectral Energy Threshold):** Spectral energy exceeding 0.7 predicts functional degradation of LoRA adapters (Biderman et al., 2024).
- **H2 (Cross-Lingual Transfer):** Initializing LoRA adapters from a related-language checkpoint (Kazakh $\rightarrow$ Kyrgyz) preserves source-language knowledge with minimal additional Frobenius perturbation.

Our contributions are as follows:

1. We provide a controlled three-way **falsification of the SE-threshold hypothesis** for low-resource quantized LoRA. Three $r=64$ runs (E5/E5b/E5c) at near-identical settled spectral trajectories span the full functional range from catastrophic collapse to baseline-better performance; SE cannot distinguish them. A full-precision BF16 control (C3) rules out quantization as the cause (Section 4.3).
2. We characterize **Frobenius-norm growth** as a useful but language-dependent diagnostic, paired with cross-lingual PPL probes: norm growth

alone fails across language categories (English $33\times$ without forgetting vs. Kazakh $38\times$ with forgetting; Section 4.4).

3. We show that **related-language warm-starts preserve source-language knowledge**: Kazakh $\rightarrow$ Kyrgyz transfer halves cross-lingual Kazakh perplexity across two seeds (23.5 ± 0.1 vs. 50.0 ± 2.4) with only $+13\%$ Frobenius growth. A cross-family control (EN $\rightarrow$ KY) rules out a generic warm-start explanation: an unrelated source yields no retention. The effect is specific to related-language pairs (Section 4.5).
4. We provide a quantitative **characterization of subword fragmentation** for Kyrgyz, Kazakh, and Uzbek and document Wikipedia-WikiANN entity overlap ($66\%/56\%/2\%$) as an upper bound on memorization-driven NER F1 (Section 3.6, Section D).
5. We release an open-source **SVD monitoring toolkit**, all training and evaluation code, and a self-contained BF16 control bundle for reproducibility.

2 Related Work

2.1 LoRA and Parameter-Efficient Fine-Tuning

Low-Rank Adaptation (LoRA) (Hu et al., 2022) parameterizes weight updates as $\Delta W = BA$ where $B \in \mathbb{R}^{d \times r}$ and $A \in \mathbb{R}^{r \times k}$ with rank $r \ll \min(d, k)$, reducing trainable parameters by orders of magnitude. QLoRA (Dettmers et al., 2023) extends this approach by quantizing the base model to 4-bit precision, enabling fine-tuning of multi-billion parameter models on consumer hardware. Subsequent work has explored rank allocation strategies (Zhang et al., 2023), training stability improvements (Lialin et al., 2023), weight-decomposed variants such as DoRA (Liu et al., 2024), alternative singular-value-based initialization (PiSSA, (Meng et al., 2024)), and theoretical analyses of low-rank update expressiveness (Malladi et al., 2023). The broader observation that LLM fine-tuning effectively operates in a low intrinsic dimension (Aghajanyan et al., 2021) provides important context for why low-rank updates work at all.

2.2 Spectral Analysis of Neural Network Weights

The spectral properties of neural network weight matrices have been studied as indicators of training dynamics and generalization. [Martin and Mahoney \(2021\)](#) demonstrated that well-trained networks exhibit heavy-tailed singular value distributions, a property they termed “implicit self-regularization.” In the context of LoRA, [Biderman et al. \(2024\)](#) showed that LoRA adapters learn less and forget less than full fine-tuning, and identified spectral collapse as a potential diagnostic for training degradation. [Jang et al. \(2024\)](#) provided a more detailed analysis of spectral collapse dynamics, proposing that monitoring the spectral energy ratio can serve as an early warning system for adapter failure. However, both studies focused primarily on English-language settings with relatively large training corpora, leaving open the question of whether these findings generalize to low-resource, morphologically complex languages.

2.3 Low-Resource NLP for Turkic Languages

Turkic languages have received increasing attention in the NLP community, though they remain severely underrepresented relative to their speaker populations. Recent benchmarks such as TUMLU ([Isbarov et al., 2024](#)) provide standardized evaluation for Turkic language understanding. Work on Kazakh NLP has leveraged Wikipedia and web-crawled corpora ([Yeshpanov et al., 2024](#)), while Kyrgyz and Uzbek remain among the least-resourced languages with significant digital presence. The challenge of tokenizer mismatch for agglutinative languages has been documented ([Rust et al., 2021](#); [Ahia et al., 2023](#)), with morphologically rich languages consistently suffering from higher subword fragmentation ratios that inflate sequence lengths and reduce the effective information density per token.

2.4 Catastrophic Forgetting in Fine-Tuning

Catastrophic forgetting ([McCloskey and Cohen, 1989](#); [Kirkpatrick et al., 2017](#))—the loss of previously learned knowledge when adapting to new tasks—is a persistent challenge in continual learning. While LoRA’s low-rank

constraint was hypothesized to mitigate forgetting by limiting the magnitude of weight updates ([Biderman et al., 2024](#)), recent work has shown that aggressive fine-tuning can still induce significant forgetting in LLMs ([Luo et al., 2023](#); [Kotha et al., 2024](#)), particularly for languages distant from the training data ([Csaki et al., 2024](#)). The cross-lingual transfer pattern we examine—warm-starting from a related-language checkpoint—is conceptually adjacent to the adapter-based cross-lingual transfer framework of MAD-X ([Pfeiffer et al., 2020](#)), although MAD-X uses dedicated language adapters rather than sequential LoRA fine-tuning. The relationship between spectral properties of adapter weights and the degree of catastrophic forgetting has not been systematically studied for multilingual LoRA in low-resource settings.

3 Experimental Setup

3.1 Base Model

We use Gemma-2-9B ([Gemma Team et al., 2024](#)) as our base model, a 9.24 billion parameter decoder-only transformer with 42 layers (per the official model configuration).¹ The model is quantized to 4-bit precision using the NF4 data type with double quantization ([Dettmers et al., 2023](#)), reducing memory requirements to approximately 5.1 billion effective parameters while preserving model quality. All experiments use bfloat16 compute precision. We chose Gemma-2-9B because it possesses baseline linguistic knowledge of Turkic languages from pretraining, allowing us to isolate the spectral dynamics of LoRA *adaptation* from the process of primary language acquisition. The functional collapse observed at $r=64$ therefore reflects a pathology of the adaptation process, not an inability of the model to represent Turkic tokens.

3.2 Training Data

We compile monolingual corpora of approximately 150 MB for each of the three target languages:

- **Kazakh (KZ):** 61,879 records sourced from Kazakh Wikipedia and the

¹<https://huggingface.co/google/gemma-2-9b/blob/main/config.json>
num_hidden_layers=42.

sozkh-corpus, totaling 14.1M tokens (151.5 MB).

- **Kyrgyz (KY):** 17,360 records from a curated local corpus comprising literature, history, and encyclopedic texts, totaling 4.4M tokens (151.4 MB).²
- **Uzbek (UZ):** 21,242 records from uz-books, Uzbek Wikipedia, and FineWeb-2 with Cyrillic script filtering, totaling 5.4M tokens (152.6 MB). Note that our Uzbek corpus uses Cyrillic script, which reflects a subset of Uzbek digital content; results may not generalize to Latin-script Uzbek, which dominates contemporary formal writing.

Additionally, we compile an **English (EN)** control corpus of 52,268 records from English Wikipedia (155 MB, 12.7M tokens), matched in size to the Turkic corpora.

All texts are chunked into segments of 200–5,000 characters. A 10% held-out split is reserved for validation. All corpora are matched by raw byte size (~ 150 MB), but under comparable subword fragmentation ($\sim 3.2\times$ tokens per word across all three languages; Section 3.6) token counts still differ substantially (KZ: 14.1M, UZ: 5.4M, KY: 4.4M). This reflects a *source-text difference*—the Kazakh corpus is dominated by dense encyclopedic prose with shorter word forms, while the Kyrgyz corpus contains a higher proportion of literary material with longer word forms. Because tokenizer fragmentation is nearly identical across the three languages, this imbalance is a property of the corpora, not of the tokenizer. We flag it explicitly as a confound for any direct E1 vs. E3 comparison (see also Limitation L5): Kazakh’s superior target PPL could partially reflect $3\times$ more training tokens rather than better tokenizer coverage for Kazakh per se. We control for this with E1b, a reduced-Kazakh subset of ~ 1.5 M training tokens—roughly $3\times$ *smaller* than the KY baseline (4.4M tokens). The original goal was a token-matched control, but the corpus subsetting (sized via $\text{word}\times\text{fragmentation-ratio}$

estimation) under-shot once actual chunked tokens were counted; the comparison still serves its purpose, as discussed in Sections 4.2 and 6.

3.3 LoRA Configuration

We apply LoRA adapters to all attention and MLP projection matrices: `q_proj`, `k_proj`, `v_proj`, `o_proj`, `gate_proj`, `up_proj`, and `down_proj` (7 modules per layer, 294 total across 42 layers). We set $\alpha = 2r$ in all experiments and use the paged AdamW optimizer with a cosine learning rate schedule and 5% warmup. Table 1 summarizes all experimental configurations.

3.4 SVD Monitoring Protocol

We perform singular value decomposition (SVD) of each `lora_B` weight matrix every 100 training steps, computing five spectral metrics:

Spectral Energy (SE) quantifies the dominance of the largest singular value:

$$\text{SE} = \frac{\sigma_1}{\sum_{i=1}^r \sigma_i} \quad (1)$$

Effective Rank (Roy and Vetterli, 2007) measures the number of significantly contributing singular values:

$$\text{EffRank} = \exp\left(-\sum_{i=1}^r p_i \log p_i\right), \quad p_i = \frac{\sigma_i}{\sum_j \sigma_j} \quad (2)$$

Stable Rank provides a continuous relaxation of matrix rank:

$$\text{StableRank} = \frac{\|W\|_F^2}{\sigma_1^2} = \frac{\sum_i \sigma_i^2}{\sigma_1^2} \quad (3)$$

We log stable rank for completeness and reproducibility but find that it is dominated by the same shape-vs-scale issue as SE in our setting: changes in stable rank are driven primarily by the denominator σ_1^2 , which itself rises during healthy training. We therefore do not analyze it further in the main text; raw values are released alongside the SVD logs.

SVD Entropy measures the uniformity of the singular value distribution:

$$H_{\text{SVD}} = -\sum_{i=1}^r p_i \log_2 p_i \quad (4)$$

²Due to licensing restrictions on the source materials, the Kyrgyz corpus cannot be publicly redistributed. Corpus statistics and preprocessing details are provided in Section A.

ID	Experiment	Language	Rank	LR	Epochs	Trainable %	Init
E1	KZ Baseline (14.1M tok)	Kazakh	16	2×10^{-4}	3	1.05%	Random
E1b	KZ Small-corpus (1.5M tok)	Kazakh	16	2×10^{-4}	3	1.05%	Random
E2	UZ Baseline	Uzbek	16	2×10^{-4}	3	1.05%	Random
E3	KY Baseline	Kyrgyz	16	2×10^{-4}	3	1.05%	Random
E4	KY Overfit	Kyrgyz	16	2×10^{-4}	10	1.05%	Random
E5	KY $r=64$ high-LR	Kyrgyz	64	5×10^{-4}	5	4.08%	Random
E5b	KY $r=64$ std-LR	Kyrgyz	64	2×10^{-4}	5	4.08%	Random
E5c	KY $r=64$ epoch-matched	Kyrgyz	64	2×10^{-4}	3	4.08%	Random
E6	KZ→KY Transfer	Kyrgyz	16	2×10^{-4}	3	1.05%	KZ checkpoint
E6b	EN→KY Transfer [◊]	Kyrgyz	16	2×10^{-4}	3	1.05%	EN checkpoint
E7	Diverged	KY [‡]	64	1×10^{-3}	3	4.08%	Random
E8	EN Control	English	16	2×10^{-4}	3	1.05%	Random
C3	KY BF16* control	Kyrgyz	16	2×10^{-4}	3	0.58%	Random

Table 1: Summary of experimental configurations. All experiments except C3 use Gemma-2-9B with 4-bit NF4 quantization. C3 uses full bfloat16 precision (no quantization) and serves as the quantization control. All runs use $\alpha = 2r$, cosine LR schedule, and effective batch size 16. [‡]E7 was trained on Kyrgyz data but diverged at step 800 due to the aggressive learning rate; it is excluded from downstream evaluation but analyzed in Section B. *C3 trainable% reads 0.58% rather than 1.05% because PEFT’s parameter count differs between BF16 and 4-bit base models; the LoRA configuration (rank, alpha, target modules) is identical to E3. [◊]E6b is a cross-family transfer control: same setup as E6 but warm-started from the unrelated English baseline (E8) instead of the related-language Kazakh baseline (E1).

Frobenius Norm captures the overall magnitude of the weight matrix:

$$\|W\|_F = \sqrt{\sum_{i=1}^r \sigma_i^2} \quad (5)$$

We monitor both `lora_A` and `lora_B` Frobenius norms separately. The spectral collapse threshold is set at $SE = 0.7$ following Biderman et al. (2024).

3.5 Evaluation Protocol

We evaluate each adapter on three tasks:

Perplexity (PPL): Computed on held-out validation splits (500 samples, 113–128K tokens per language). We construct a *cross-lingual perplexity matrix* by evaluating each adapter on all three languages, providing a direct measure of catastrophic forgetting.

Named Entity Recognition (NER): Using the WikiANN dataset (Pan et al., 2017) with 3-shot prompting on 100 test samples per language. We report two complementary metrics: (i) generation-based F1 micro-averaged across PER/ORG/LOC types (subject to parse-failure noise and ± 0.085 Wilson CI at $n=100$); and (ii) log-likelihood span typing accuracy on gold-span boundaries (Section C)—a parse-failure-immune probe that scores entity-type knowledge directly. The

log-likelihood metric is our primary NER number in Table 3; F1 is reported alongside for comparability with prior work. **Caveat for Uzbek:** our UZ corpus uses Cyrillic script while WikiANN-uz is primarily Latin (overlap 2%, see L9 and Section D); cross-configuration UZ NER comparisons should therefore be read as cross-script transfer rather than in-distribution NER.

Question Answering (QA): Using TUMLU (Isbarov et al., 2024), a multiple-choice benchmark for Turkic language understanding, with 5-shot log-likelihood evaluation on 695–794 questions per language.

3.6 Subword Tokenization Analysis

We measure subword fragmentation as the average number of tokens per whitespace-delimited word for each language. As shown in Table 2, all three Turkic languages exhibit severe fragmentation at approximately $3.2\times$ the rate of English, reflecting the fundamental mismatch between the Gemma tokenizer (optimized for English and high-resource languages) and agglutinative Turkic morphology.

4 Results

4.1 Baseline Performance

Table 3 presents the primary evaluation metrics across all experiments. Several patterns

Language	Tokens/Word	Ratio to EN
English	1.15	1.00×
Uzbek	3.68	3.20×
Kyrgyz	3.69	3.21×
Kazakh	3.74	3.25×

Table 2: Subword fragmentation ratios for the Gemma-2-9B tokenizer. Turkic languages require approximately $3.2\times$ more tokens per word than English, reflecting severe tokenizer mismatch for agglutinative morphology.

emerge from the baseline comparisons.

The Kazakh baseline (E1) achieves the best target-language perplexity (2.73), likely reflecting both the larger effective token count (14.1M vs. 4.4–5.4M) and better representation in Gemma’s pretraining data. The Uzbek baseline (E2) achieves the highest NER F1 on its own language (0.287), though its perplexity (4.03) is higher than Kazakh. The Kyrgyz baseline (E3) shows the weakest overall performance (PPL 4.78, NER F1 0.157), consistent with Kyrgyz being the least resourced of the three languages.

We note apparent cross-lingual NER anomalies: the KZ baseline (E1) achieves higher NER F1 on Uzbek (0.578) than on Kazakh itself (0.164). Wilson confidence intervals at $n=100$ are wide (e.g., E6 KY NER: 0.253 ± 0.085 , 95% CI [0.168, 0.339]), and the E3 vs. E6 comparison (0.157 vs. 0.253) shows overlapping intervals. We therefore interpret absolute NER F1 values as noisy point estimates and focus our analysis on: (1) the qualitative distinction between functional models ($F1 > 0$) and collapsed ones ($F1 = 0$, which is unambiguous), and (2) consistent directional patterns across multiple metrics (PPL, NER, TUMLU) rather than any single number.

The most striking result is the behavior of the $r=64$, high-LR configuration (E5): despite achieving a target-language perplexity of only 5.90 (a modest 23% increase over the baseline), it suffers *complete* NER failure with $F1 = 0$ across all three languages, driven by 100% parse failures, and its TUMLU accuracy drops to 23.2%—near the 25% random baseline for 4-option multiple choice. This indicates that the model has lost both instruction-following and latent language-understanding capability,

a form of functional collapse that perplexity alone fails to capture.

Disentangling rank, learning rate, and training duration (E5 vs. E5b vs. E5c). The original E5 ($r=64$, $lr=5\times 10^{-4}$, 5 epochs) varies three factors at once relative to the E3 baseline. We isolate each in turn. **E5b** matches rank but lowers LR to 2×10^{-4} (replicated at seeds 42 and 123; Section F): catastrophic collapse vanishes (NER F1 ≥ 0.12 , TUMLU 33.2%, KY PPL 4.55–4.62, all far outside E5’s collapsed regime), identifying *learning rate* as the proximate cause of collapse. **E5c** additionally matches the epoch budget to baseline (3 epochs) and reveals an even cleaner picture: at matched rank, LR, and epochs, $r=64$ *outperforms* the $r=16$ baseline on the target language (KY PPL 3.86 vs. 4.78; TUMLU 38.5% vs. 34.7%). Frobenius growth at matched epochs is also virtually identical to baseline ($8.7\times$ vs. $8.4\times$; Table 5), demonstrating that *rank itself contributes capacity, not pathology*. The cross-lingual cost is the residue: E5c still loses more KZ knowledge than the $r=16$ baseline (KZ PPL 93.5 vs. 47.6), so higher rank does increase forgetting under matched conditions, but at $r=64$ this is a mild $2\times$ penalty rather than the $14\times$ catastrophe of E5. The full causal decomposition: catastrophic collapse = aggressive LR; cross-lingual forgetting growth = training duration; rank-only effect = better target-language fit at $\sim 2\times$ baseline cross-lingual penalty. Importantly, E5, E5b, and E5c all show mean SE well below 0.1 at every checkpoint, with final max SE between 0.06 and 0.12 (compared to a brief early-training peak around 0.21–0.25 that decays in the first ~ 100 steps), yet they span the full range from catastrophic collapse (E5) to baseline-better target performance (E5c)—another demonstration (complementing the EN control in Section 4.4) that spectral metrics do not discriminate these regimes.

Entity knowledge survives in E5 (log-likelihood NER). To separate instruction-following loss from genuine knowledge loss, we additionally score each gold WikiANN span by log-likelihood over the three type labels (Section C). Under this parse-failure-immune evaluation, E5 achieves the *highest* target-language (KY) type accuracy in our sweep

ID	Experiment	Target Lang	PPL (target)	NER TypeAcc (KY)	NER F1 (gen) (KY, ± 0.09 CI)	TUMLU (target)
E1	KZ Baseline (14.1M tok)	KZ	2.73	57.7%	0.219	38.4%
E1b	KZ small-corpus (1.5M tok)	KZ	4.17	46.0%	0.150	35.6%
E2	UZ Baseline	UZ	4.03	54.9%	0.138	35.4%
E3	KY Baseline	KY	4.78	50.4%	0.157	34.7%
E4	KY Overfit (10ep)	KY	4.18	56.8%	0.173	33.9%
E5	KY $r=64$ high-LR	KY	5.90	59.5% [‡]	0.000	23.2%
E5b	KY $r=64$ std-LR (5ep)	KY	4.55	53.1%	0.146	33.2%
E5c	KY $r=64$ epoch-matched (3ep)	KY	3.86	54.9%	0.115	38.5%
E6	KZ→KY	KY	4.73	59.5%	0.253	33.7%
E6b	EN→KY (cross-family)	KY	4.78	54.0%	0.133	36.1%
E8	EN Control	EN	5.54 [*]	56.8%	0.209	37.7% [†]
C3	KY BF16 (no quant)	KY	4.99	58.6%	0.164	36.1%

Table 3: Main evaluation results across all experiments. **NER TypeAcc (KY)** is our primary NER metric—log-likelihood span typing on gold WikiANN-ky boundaries (Section C); it is parse-failure immune and has lower seed-variance than generation F1. **NER F1 (gen)** is the generation-based score with ± 0.085 Wilson 95% CI at $n=100$, retained for comparability with prior work. We report cross-lingual NER and KZ TUMLU in the appendix (Sections G and H). **E5 vs. E5b/E5c disentangles the rank-LR-epoch confound:** catastrophic NER parse failure at high LR (E5, F1=0, TUMLU near random) vanishes at baseline LR (E5b/E5c) and at matched epochs (E5c) actually *outperforms* the $r=16$ baseline. ^{*}EN PPL on Kazakh; [†]EN TUMLU on Kyrgyz. [‡]E5’s high TypeAcc despite F1=0 shows entity knowledge survives the collapse—it is structured-output generation that fails (Section C).

(59.5%, tied with the transfer run E6), despite scoring F1 = 0 under generation. This confirms that E5’s collapse is specifically a loss of structured-output instruction-following, not an erasure of entity-type knowledge. Cross-lingual degradation still appears—E5 drops to 50.0% on Kazakh versus 70–73% for the $r=16$ baselines—so the higher-rank, higher-LR combination does erode non-target-language representations, just not to the extent that the generation-based F1 = 0 suggests. Relatedly, the transfer advantage that was +61% relative under generation F1 becomes +9.1 absolute percentage points under log-likelihood (59.5% vs. 50.4%), which we consider a more honest directional effect-size estimate.

4.2 Cross-Lingual Degradation

The cross-lingual perplexity matrix (Table 4) reveals the extent of catastrophic forgetting induced by monolingual fine-tuning.

Several key observations emerge. First, the $r=64$ configuration (E5) induces catastrophic forgetting: Kazakh PPL rises to 659.25 and Uzbek to 742.73, representing increases of over $240\times$ and $180\times$ relative to the respective baselines. Second, the KZ→KY transfer (E6) retains substantially more Kazakh knowledge than direct Kyrgyz training: KZ PPL

of 23.49 versus 47.58 for the KY baseline (a $2\times$ improvement). Third, extended training (10 epochs, E4) nearly doubles cross-lingual PPL relative to the 3-epoch baseline (87.65 vs. 47.58 for KZ; 95.76 vs. 56.85 for UZ), confirming that additional training steps systematically erode non-target-language knowledge. Fourth, the Uzbek baseline (E2) shows the highest asymmetric forgetting: KY PPL rises to 116.89, substantially worse than the reverse direction (UZ PPL under KY training is 56.85), possibly reflecting the greater script distance between Uzbek (historically Latin) and Kyrgyz (Cyrillic).

Data-volume vs. language-coverage contributions (E1 vs. E1b). To disentangle whether E1’s strong target PPL (2.73) reflects its $3\times$ larger token count or better tokenizer/pretraining coverage for Kazakh, we ran a small-corpus control E1b on a Kazakh subset of ~ 1.5 M training tokens (intended as a token-matched control; the actual chunked-token count under-shot the 4.4M target). The target KZ PPL rises to 4.17—showing that data volume contributes a meaningful ~ 1.4 PPL improvement to E1 (gained over $\sim 9\times$ more training tokens). Critically, E1b’s 4.17 remains *below* the Kyrgyz baseline’s 4.78 even though E1b trained on $\sim 3\times$ fewer to-

Trained on	KY PPL	KZ PPL	UZ PPL
KZ Baseline (14.1M tok)	40.75	2.73	44.27
KZ Small-corpus (1.5M tok, E1b)	25.55	4.17	33.17
UZ Baseline	116.89	64.44	4.03
KY Baseline	4.78	47.58	36.85
KY Overfit	4.18	87.65	55.76
KY $r=64$ (lr 5×10^{-4} , 5ep, E5)	5.90	659.25	143.33
KY $r=64$ (lr 2×10^{-4} , 5ep, E5b)	4.55	128.04	162.68
KY $r=64$ (lr 2×10^{-4} , 3ep, E5c)	3.86	93.52	111.25
KZ→KY (related)	4.73	23.49	24.19
EN→KY (cross-family, E6b)	4.78	47.46	63.42
EN Control	14.96	5.54	16.55
KY BF16 (C3, no quant)	4.99	43.97	49.02

Table 4: Cross-lingual perplexity matrix. Each row is an adapter; each column is the evaluation language. Bold indicates target-language PPL. The high-LR $r=64$ configuration (E5) shows catastrophic cross-lingual degradation (> 600 KZ PPL), but the controlled-LR $r=64$ (E5b, lr= 2×10^{-4}) reduces this 5-fold to 128 while improving target KY PPL to 4.55—isolating learning rate, not rank, as the driver of catastrophic collapse. Higher rank alone still roughly doubles cross-lingual forgetting versus the $r=16$ KY baseline (128 vs. 47 KZ PPL). The EN control shows minimal Turkic forgetting despite comparable Frobenius norm growth ($33\times$).

kens than the KY baseline (1.5M vs. 4.4M); Kazakh therefore enjoys a genuine pretraining-coverage advantage that is not attributable to corpus size. The two factors are complementary rather than reducible. Cross-lingually, E1b shows substantially lower forgetting than E1 (KY PPL 25.55 vs. 40.75; UZ PPL 23.17 vs. 41.27), consistent with fewer training steps inducing milder weight perturbations—a finding that parallels the extended-training penalty observed for E4 in the overfit paradox discussion.

Critically, the **EN control experiment** (E8) reveals that Frobenius norm growth alone does not predict forgetting. Despite exhibiting $33\times$ norm growth—comparable to the KZ baseline ($38\times$)—the EN adapter induces minimal Turkic language degradation (KZ PPL 5.54 vs. 47.58 for the KY baseline). This is consistent with forgetting severity depending not only on the magnitude of weight perturbations but also on their *relationship* to pre-trained representations. We hypothesize that English weight updates largely reinforce already well-represented features (operating in directions the base model already encodes), while Turkic updates must instead restruct-

ure unrepresented regions of the weight space, causing disproportionate interference with other languages. We do not directly measure update direction (e.g., via cosine similarity with pre-trained weights) in this work and note this as a future-work avenue (Section 6). Figure 1 visualizes these patterns as a heatmap.

4.3 Spectral Dynamics: Challenging H1

Our central finding is that the spectral energy threshold of 0.7 proposed by Biderman et al. (2024) is **never reached** in any of our experiments, including those exhibiting severe functional degradation. Figure 2 shows the evolution of mean and maximum SE across training steps for representative experiments.

The KY baseline ($r=16$) shows maximum SE peaking near 0.44 in the first ~ 100 steps and then *decreasing* to ~ 0.32 at the final checkpoint, well below the 0.7 threshold throughout. Both $r=64$ configurations—the collapsed E5 and the healthy-adapting E5b—show the same qualitative pattern but at much smaller values, ending at final mean SE ~ 0.03 and final max SE ~ 0.06 – 0.11 (well below the threshold). The fact that E5 and E5b end with near-identical spectral profiles yet completely different functional behavior (NER F1 = 0 vs. 0.146; TUMLU 23% vs. 33%) is a direct falsification of SE as a diagnostic in our setting: SE is a *shape* metric that cannot distinguish the two regimes. As training progresses, all singular values grow, but the growth is distributed across multiple directions, causing the relative contribution of σ_1 to decrease—regardless of whether the adapter has collapsed functionally.

The effective rank tells a complementary story: for both $r=64$ configurations, effective rank *increases* from initialization, indicating that the adapter is utilizing more of its available capacity. Under standard interpretations, increasing effective rank signals healthy training dynamics. Yet the high-LR E5 configuration suffers complete NER failure and catastrophic cross-lingual forgetting despite this “healthy” trajectory, while the controlled-LR E5b does not—again showing that the spectral trajectory alone cannot distinguish the two.




Figure 1: Cross-lingual perplexity heatmap. Darker cells indicate higher perplexity (more forgetting). The $r=64$ row is clearly pathological, while the KZ→KY transfer row shows the most balanced profile.

Figure 3 shows the layer-wise SE distribution at the final checkpoint, confirming that no individual layer approaches the collapse threshold.

Quantization is not the cause (BF16 control, C3). A reviewer-prompted concern was that 4-bit NF4 quantization might itself suppress spectral growth and account for the sub-threshold SE. We address this with a dedicated full-precision (bfloat16, no quantization) replication of the KY baseline (E3 setup, $r=16$, $lr=2\times 10^{-4}$, 3 epochs) on a cloud A100. The result confirms, rather than weakens, the SE finding: BF16 and 4-bit produce essentially indistinguishable SE behaviour. In BF16 the final-checkpoint max SE is 0.32 and final mean SE is 0.10, matching 4-bit E3 within seed variance (0.33 ± 0.01 and 0.10 across two seeds; Table 11); the in-training peak max SE is also comparable (0.47 in BF16 vs. 0.45 ± 0.02 in 4-bit E3 across seeds). Frobenius norm growth at the same point is $8.6\times$ in BF16 versus $8.4\times$ in E3—indistinguishable. Cross-lingual perplexity, NER F1, log-likelihood span typing, and TUMLU accuracy are all within ± 1 standard deviation of the seed-level variance we observe for E3 across two seeds (Section F and table 11). Quantization granularity therefore does not materially alter spectral or downstream behaviour in our setting; the SE-

threshold falsification carries over from 4-bit to full precision.

H1 is not supported in our setting: spectral energy below the proposed threshold does *not* guarantee healthy training, and SE can decrease during training that leads to functional collapse. We note that our experiments use 4-bit NF4 quantization, which may itself suppress singular value growth relative to the full-precision setting of Biderman et al. (2024); our findings therefore challenge the generality of the SE threshold rather than refuting it unconditionally. The SE metric fundamentally measures the *shape* of the singular value distribution while ignoring its *scale*—a critical limitation when diagnosing pathological training regimes in quantized, low-resource settings.

4.4 Frobenius Norm: The Missing Diagnostic

If spectral energy fails as a diagnostic, what about scale-aware alternatives? We examine the **Frobenius norm growth ratio** of the `lora_B` matrices, which captures the magnitude of weight perturbations that SE discards. While Frobenius growth provides better separation than SE in some regimes, our cross-language and cross-rank analysis shows it is also *not* a universal diagnostic.

Table 5 and Figure 4 reveal a nuanced pic-

SVD Dynamics: Baseline vs Aggressive vs Transfer




Figure 2: Spectral energy dynamics during training. Top: mean SE across all layers. Bottom: maximum SE across layers. No experiment approaches the proposed 0.7 threshold. The $r=64$ configuration shows *decreasing* SE despite suffering the worst functional degradation.




Figure 3: Layer-wise spectral energy at the final training checkpoint. Each row is an experiment; each column is a transformer layer. No layer in any experiment exceeds the 0.7 threshold.

ture. The KZ→KY transfer exhibits only +13% Frobenius norm growth, while the KY baseline shows $\sim 8\times$ growth and the pathological $r=64$ configuration shows 15–38 $\times$ growth. However, the EN control experiment complicates this gradient: despite $\sim 33\times$ norm growth—comparable to the KZ baseline (38 $\times$)—EN training causes minimal cross-lingual forgetting (KZ PPL 5.52 ± 0.02 across 2 seeds vs. ~ 48 for the KY baseline with only 8 $\times$ growth). This EN-vs-Turkic dissociation is consistent across two seeds (E8 Frobenius growth 33.39 $\times$ at seed 42 and 33.17 $\times$ at seed 123; see Section F). It reveals that Frobenius norm growth is a reliable diagnostic *within* a language category (underrepresented vs. well-represented) but not *across* categories: the same magnitude of weight perturbation has qualitatively different effects depending on

how well the target language is represented in the pretrained model.

Figure 5 visualizes the relationship between Frobenius norm growth and cross-lingual PPL degradation. We do not perform a formal statistical test, as our sample size ($n=7$ Turkic experiments) is far too small for a meaningful rank-correlation test. Visually, within the Turkic experiments the trend is in the expected direction (more growth, more forgetting), but the EN control is a clear outlier (33 $\times$ growth with minimal forgetting), indicating that norm growth alone is not sufficient to predict catastrophic forgetting. The Frobenius norm is a *scale-aware* metric that captures the total magnitude of weight updates, but the *direction* of those updates—whether they reinforce or disrupt pretrained representations—is plausibly what determines

Frobenius Norm Dynamics: The Missing Diagnostic




Figure 4: Frobenius norm dynamics of `lora_B` matrices during training. The $r=64$ configuration shows dramatically larger norm growth compared to $r=16$ baselines. The transfer experiment begins at the KZ baseline’s final norm and shows minimal additional growth.

Experiment	$\ B\ _F^{\text{init}}$	$\ B\ _F^{\text{final}}$	Metric	KY Baseline	KZ→KY
KZ→KY Transfer (related)	3.043	3.436	PPL (KY)	4.78	4.73
EN→KY Transfer (cross-family, E6b)	0.653	2.023	PPL (KZ)	47.58	23.49
KY Baseline $r=16$	0.227	1.899	PPL (UZ)	56.85	124.19
EN Control $r=16$	0.092	3.076	NER F1 (KY)	0.157	0.253
KY BF16 $r=16$ (C3)	0.218	1.880	NER F1 (KZ)	0.190	0.130
KZ Baseline $r=16$	0.081	3.107	NER F1 (UZ)	0.326	0.348
KY $r=64$ lr 5×10^{-4} , 5ep (E5)	0.566	8.859	TUMLU (KY)	34.7%	33.7%
KY $r=64$ lr 2×10^{-4} , 5ep (E5b)	0.247	4.486	TUMLU (KZ)	30.2%	31.7%
KY $r=64$ lr 2×10^{-4} , 3ep (E5c)	0.388	3.385	$\ B\ _F$ growth	8.36×	1.13×
Diverged $r=64$	—	—	explodes		

Table 5: Frobenius norm growth ratios for `lora_B` matrices. EN control and KZ baseline show comparable growth (~ 33 – $38\times$), yet EN produces minimal cross-lingual forgetting (KZ PPL 5.54) while KY baseline with only $8\times$ growth causes severe forgetting (KZ PPL 47.58). This demonstrates that norm growth alone is insufficient as a universal diagnostic—language representation in the pretrained model modulates the impact of weight perturbations. Note also that E5 and E5b exhibit comparable Frobenius growth ($15.6\times$ vs. $18.2\times$), yet only E5 suffers NER parse failure and near-random TUMLU: norm growth does not discriminate between the catastrophic-collapse and moderate-forgetting regimes when rank is held constant.

downstream impact (a hypothesis we do not directly verify; see Section 6).

4.5 Cross-Lingual Transfer

The KZ→KY transfer experiment (E6) initializes LoRA adapters from the final Kazakh checkpoint (E1) before continuing training on Kyrgyz data. This simple sequential transfer yields substantial improvements across multi-

Table 6: Comparison of direct Kyrgyz training (E3) and Kazakh→Kyrgyz transfer (E6), reported at seed 42. Transfer’s most robust effect is KZ retention (PPL 23.49 vs. 47.58; 23.5 ± 0.1 vs. 50.0 ± 2.4 across two seeds—Section F). The single-seed KY NER F1 improvement of +61% (0.253 vs. 0.157) is *not* robust to reseeding (seed 123: 0.179 vs. 0.182); we therefore characterize transfer as a retention mechanism rather than a target-NER improvement. UZ PPL increases under transfer (124.19 vs. 56.85), suggesting the benefit is specific to closely related Kipchak languages.

ple metrics (Table 6).

The most visible improvement at seed 42 is Kyrgyz NER F1: 0.253 versus 0.157, a relative improvement of +61%. However, a second seed (123) gives 0.179 vs. 0.182 for the same comparison—the NER-F1 advantage is not robust across seeds (see Section F). Target-language perplexity is essentially identical under both seeds (4.73–4.78 for E6 vs. 4.78–4.83 for E3), indicating that the transfer does not sacrifice language modeling quality but also that it does not noticeably improve it. The one *robust* transfer effect is Kazakh retention:




Figure 5: Frobenius norm growth ratio versus cross-lingual perplexity degradation. Larger norm growth is associated with more severe catastrophic forgetting, providing a practical early-stopping criterion.

KZ PPL of 23.5 ± 0.1 versus 50.0 ± 2.4 across seeds, a $2\times$ improvement well outside seed variance, with KZ TUMLU accuracy improving consistently (31.7–32.6% vs. 30.2–32.6% at seed 42 and 123 respectively).

Interestingly, even Uzbek NER F1 improves slightly under transfer (0.348 vs. 0.326), suggesting some positive transfer to the third Turkic language. However, Uzbek PPL degrades substantially (124.19 vs. 56.85, a $2.2\times$ increase). This asymmetry has practical implications: sequential transfer between closely related languages (Kazakh and Kyrgyz are both Kipchak-branch, Cyrillic-script) can *harm* more distant relatives (Uzbek, Karluk-branch, historically Latin-script). The adapter, having been optimized first for Kazakh then for Kyrgyz, specializes its weight updates in the Kipchak direction, displacing representations useful for Uzbek. Practitioners considering chain-of-transfer strategies should therefore monitor not only source-target retention but also collateral damage to typologically distant languages.

Cross-family control (E6b): is any warm-start enough? A natural concern with the above interpretation is that the KZ-retention effect in E6 could simply reflect “the adapter was trained on KZ first, so of course it remembers KZ”—in which case any warm-

started adapter, even from an unrelated language, would yield similar retention of its own source. We address this with a cross-family control (E6b): same setup as E6 but warm-started from the English baseline (E8) instead of the Kazakh baseline (E1). If *any* warm-start is sufficient, then EN→KY should preserve the EN-baseline cross-lingual profile in the same way KZ→KY preserves the KZ profile.

The data falsify this null. EN→KY converges to KZ PPL 47.46—essentially identical to direct KY training (47.58) and far above E6’s 23.49. The KY target and TUMLU values for EN→KY also align with direct training rather than with E6. Frobenius growth tells a consistent story: KZ→KY needs only $1.13\times$ growth, EN→KY needs $3.10\times$, and direct training needs $8.36\times$, suggesting that the EN-initialized adapter is geometrically further from a good KY solution than the KZ-initialized one.

H2 is supported only for related-language warm-starts: cross-lingual transfer preserves source-language knowledge specifically when the source and target are closely related (here, both Kipchak, both Cyrillic). The mechanism appears to be *better initialization*: the KZ-trained adapter is already close to a good KY solution, so the additional perturbation needed is small

(+13% Frobenius growth vs. $8.36\times$ for direct training), and this low-perturbation regime is what preserves the unrelated knowledge encoded in the warm-start. An EN-warm-start, which sits further from a useful KY solution, requires a larger perturbation that effectively overwrites the source. Related-language pretraining thus acts as a kind of geometric regularizer; this is *not* a generic warm-start effect.

4.6 The Overfit Paradox

The 10-epoch Kyrgyz experiment (E4) reveals a paradoxical pattern (Figure 6): extended training *improves* target-language perplexity (4.18 vs. 4.78 for the 3-epoch baseline) despite clear indicators of overfitting (final-step train loss ~ 0.17 vs. ~ 1.46 for the 3-epoch baseline at end of training, indicating severe memorization).

However, this improvement comes at a steep cost in cross-lingual knowledge. Kazakh PPL nearly doubles (87.65 vs. 47.58), Uzbek PPL increases by 68% (95.76 vs. 56.85), and TUMLU accuracy on Kazakh drops from 30.2% to 27.8%. The spectral profile remains healthy throughout all 10 epochs ($SE < 0.5$), further confirming that SE is not a reliable indicator of training pathology.

This pattern suggests a regime of “useful overfitting” for the target language: the model memorizes patterns specific to Kyrgyz text more deeply, yielding better held-out perplexity on in-distribution data, but at the expense of broader multilingual competence. For practitioners focused on a single target language, extended training may be acceptable; for those seeking to maintain multilingual capability, it represents a significant regression.

5 Analysis and Discussion

5.1 Why Spectral Energy Fails as a Diagnostic

The failure of the SE threshold can be understood through the mathematics of LoRA initialization. Standard LoRA initializes A with Kaiming uniform and B with zeros. After the first gradient update, B acquires small random-scale singular values with roughly uniform magnitudes, yielding $SE \approx 1/r$. As training progresses, all singular values grow, but

the growth is typically distributed across multiple directions, causing SE to remain stable or even *decrease*.

For higher-rank configurations ($r=64$), the SE denominator includes more singular values, mechanically driving SE to lower values ($1/64 \approx 0.016$ at initialization versus $1/16 = 0.0625$ for $r=16$). This reveals SE as a *rank-biased metric*: since $SE \approx 1/r$ at initialization, higher-rank adapters *appear* spectrally healthier by definition, creating a dangerous trap for practitioners who monitor SE across configurations with different ranks. A seemingly healthy low-SE value for $r=64$ may actually correspond to a more pathological state than a higher SE for $r=16$.

The fundamental limitation is that SE is a *shape* metric, invariant to rescaling of the singular value spectrum. A weight matrix with singular values $[100, 50, 25]$ and one with $[0.001, 0.0005, 0.00025]$ have identical SE, despite representing vastly different magnitudes of weight perturbation. In the low-resource setting, where models must substantially restructure their representations to accommodate poorly-covered languages, the *scale* of these perturbations matters critically.

Figure 7 visualizes the singular value distributions for representative experiments, illustrating how the $r=64$ configuration develops a well-distributed but extremely large-magnitude spectrum that SE fails to flag.

5.2 Frobenius Norm as a Scale-Aware Alternative

The Frobenius norm, $\|W\|_F = \sqrt{\sum_i \sigma_i^2}$, naturally captures the scale of weight perturbations that SE discards. Its squared value equals the sum of squared singular values, providing a direct measure of the total energy in the weight update. The Frobenius norm growth ratio (final/initial) thus quantifies how much the adapter has perturbed the pretrained weights.

Rather than prescribing fixed thresholds, we report the observed growth ranges in our experiments and caution that any such numbers are derived *post hoc* from the same data used to define them, requiring held-out validation before deployment as decision rules:

- **Transfer initialization** (E6): $1.13\times$ —negligible additional perturbation atop an



Figure 6: The overfit paradox: 10 epochs of Kyrgyz training (E4) yields better target-language PPL (4.18 vs. 4.78) despite increasing cross-lingual degradation. Train loss continues to decrease while cross-lingual PPL rises monotonically.




Figure 7: Singular value distributions at the final checkpoint. The $r=64$ configuration (right) shows a well-distributed spectrum (low SE) with dramatically larger absolute magnitudes compared to the $r=16$ baseline (left), explaining why SE fails to detect pathological training.

already-adapted checkpoint.

- $r=16$ **baselines on Turkic** (E3): $8\times$, with significant cross-lingual forgetting.
- $r=16$ **on well-represented language** (E8, EN): $33\times$, with minimal cross-lingual forgetting.
- $r=16$ **on higher-resource Turkic** (E1, KZ): $38\times$, with moderate forgetting.
- $r=64$, **high LR** (E5): $15\text{--}38\times$ over longer training, with complete functional collapse.
- **Diverged training** (E7): unbounded exponential growth, reaching $19\times$ in just 800 steps before divergence.

These growth ratios refer to $\|B\|_F^{\text{final}} / \|B\|_F^{\text{init}}$ and are normalized across initialization schemes (standard LoRA, PiSSA (Meng et al., 2024), or transfer), making the relative measure more portable than absolute norms.

Why we do not propose fixed numeric thresholds. The $33\times$ versus $38\times$ contrast between EN and KZ baselines—same magnitude, opposite forgetting outcomes—is the central caution: a single threshold applied across language categories would misclassify at least one of these regimes. Rather than a pass/fail number, we suggest practitioners (i) track the *trajectory* of $\|B\|_F$ growth across training rather than a single endpoint value, (ii) flag *super-linear* growth patterns (particularly exponential profiles, as in E7) as early

warnings of divergence independent of magnitude, and (iii) always pair norm monitoring with a cross-lingual perplexity probe, since norm growth alone cannot resolve the well-represented vs. underrepresented asymmetry we document. Converting these observations into validated thresholds requires replication on held-out configurations and additional language families, which we leave to future work.

5.3 Role of Subword Fragmentation

The $3.2\times$ subword fragmentation ratio observed for all three Turkic languages has several compounding effects on LoRA adaptation:

Sequence length inflation: A Kyrgyz sentence of 50 words requires approximately $50 \times 3.69 = 185$ tokens, versus $50 \times 1.15 = 58$ tokens for the English equivalent. With a maximum sequence length of 256 tokens, each training example contains substantially less textual content, reducing the information density of each gradient update.

Morphological fragmentation: Agglutinative suffixes that carry grammatical meaning are split across multiple subwords, forcing the model to reconstruct morphological relationships across token boundaries. For instance, a single Kazakh word form may be split into 4–5 subwords, with the root, derivational suffixes, and inflectional endings each occupying separate tokens.

Link to forgetting severity (interpretive). Our EN control experiment (E8) shows a striking dissociation between norm growth and forgetting: comparable Frobenius growth (EN: $33\times$ vs. KZ: $38\times$) but minimal vs. severe cross-lingual degradation. We *hypothesize* that this dissociation reflects directionality of the weight updates—Turkic fine-tuning may need to restructure representations in regions poorly covered by pretraining, while EN fine-tuning may reinforce already well-represented features—but our experiments do not directly measure update direction. Establishing this causally would require direct geometric probes (cosine similarity between ΔW and pretrained W , or projection onto language-specific subspaces); we leave this to future work. Subword fragmentation provides a plausible mechanism for why Turkic updates would be more disruptive (more tokens per word means more constituent representations to modify), but again

we do not isolate this from other factors.

5.4 Practical Recommendations

Based on our findings, we offer the following recommendations for practitioners fine-tuning LLMs on low-resource, morphologically complex languages:

1. **Monitor Frobenius norm growth trajectories**, not spectral energy. Track the growth profile across training (not a single endpoint) and flag super-linear or exponential patterns as early-warning signals of divergence; we deliberately do not prescribe fixed numeric cut-offs, since our observations (Section 4.4) show the same growth ratio ($33\text{--}38\times$) can coexist with or without forgetting depending on language representedness.
2. **Avoid aggressive rank–LR combinations** for small corpora (~ 150 MB). The combination of $r=64$ with $\text{lr} \geq 5 \times 10^{-4}$ leads to functional collapse in our setting, even when standard spectral metrics appear healthy. Note that E5 confounds rank and learning rate; a controlled $r=64$, $\text{lr}=2 \times 10^{-4}$ experiment would be needed to isolate the rank effect.
3. **Use related-language transfer** when available. Initializing from a related-language checkpoint provides better weight initialization and substantially improves downstream task performance with minimal additional Frobenius norm growth.
4. **Apply cross-lingual perplexity monitoring** as a forgetting diagnostic. Target-language PPL alone is insufficient; catastrophic forgetting manifests primarily in non-target languages.
5. **Consider early stopping by cross-lingual PPL**, not training loss. Extended training may improve target-language metrics while severely degrading multilingual capability.

6 Limitations

Our study has several important limitations, which we discuss here in decreasing order of severity for downstream interpretation.

(L1) Limited-seed experiments. Most configurations in Table 1 are run once ($n=1$). We provide $n=2$ seed-replicated runs (seeds 42 and 123) for the seven configurations underlying our main quantitative claims: E1b (small-corpus KZ), E3 (KY baseline), E5 (catastrophic collapse), E5b (rank-LR control), E5c (epoch-matched rank control), E6 (KZ→KY transfer), and E8 (EN control); the full comparison is in Section F. The two-seed data confirm five claims that are consistent across both seeds: (a) the $2\times$ cross-lingual Kazakh-PPL retention effect of transfer (50.0 ± 2.4 direct vs. 23.5 ± 0.1 transfer); (b) E5c outperforms E3 at matched conditions in both seeds (KY PPL 3.89 ± 0.03 vs. 4.81 ± 0.03); (c) the rank-LR disentanglement (E5b in both seeds outside E5’s collapsed regime); (d) the E1b data-vs-coverage decomposition (KZ PPL 4.16 ± 0.01); and (e) the EN-vs-Turkic Frobenius dissociation (E8 $33.3 \pm 0.1\times$ growth with KZ PPL 5.52 ± 0.02 across seeds, vs. KZ baseline’s $38.5\times$ with severe forgetting). With $n=2$ we report differences as consistent-across-two-seeds rather than as population-variance estimates; the absence of evidence for a sampling-distribution claim is not evidence of robustness, and we caution readers accordingly. **Critically, the cross-family control E6b (EN→KY) is $n=1$.** The single-seed NER-F1 transfer headline (+61%) does not survive re-seeding for E6, which is why we reframe its benefit as *source-language retention*; we cannot rule out that the KZ-retention contrast between E6 and E6b could shrink under re-seeding of E6b. A second-seed E6b run is the most important remaining follow-up; replication of E1, E2, E4, and the BF16 control C3 at $n\geq 3$ would also tighten these numbers.

(L2) Training-set contamination for NER (quantified). Our Turkic training corpora include Wikipedia dumps, and the WikiANN evaluation set is itself derived from Wikipedia entity links (Pan et al., 2017). To quantify this overlap, we computed the fraction of unique gold WikiANN-test entity strings that appear verbatim somewhere in each training corpus (case-insensitive substring match): **Kazakh 65.7% (497/757), Kyrgyz 55.9% (52/93), Uzbek 2.2% (13/602).** The Kazakh and Kyrgyz NER F1 numbers in Table 3 therefore upper-bound the

contribution of memorization at $\sim 60\%$, and cross-configuration NER comparisons within those two languages are partially confounded by this shared-entity floor. The Uzbek overlap of 2%, by contrast, is below the resolution of memorization (this turns out to be a script-mismatch artifact—see L9—which makes Uzbek a clean cross-lingual probe). The transfer-induced UZ improvements in Table 3 (e.g., E1’s 0.578 F1 on UZ) are thus most likely *not* memorization. A non-Wikipedia NER benchmark such as KazNERD (Yeshpanov et al., 2022) (gated; access pending) would tighten this further; the per-type breakdown is reported in Section D.

(L3) Rank-LR-epoch confound in E5 (fully decomposed by E5b and E5c). Experiment E5 ($r=64$, $lr=5 \times 10^{-4}$, 5 epochs) originally varied three factors simultaneously. We add two controls: E5b matches rank to E5 but lowers LR; E5c additionally matches epoch count to the baseline. Together they decompose the variation cleanly (Section 4.1): aggressive LR causes catastrophic collapse, extra epochs increase cross-lingual forgetting, and rank itself (at matched LR and epochs) actually improves target-language fit while incurring only a mild cross-lingual penalty. The remaining confound is initialization: E5b/E5c keep $\alpha/r = 2$ as in E5 (giving $\alpha = 128$ at $r = 64$) versus $\alpha = 32$ at the $r = 16$ baselines, so the LR effective at the LoRA weights differs across rank conditions. A follow-up holding the effective LR at adapter weights constant ($\alpha = 32$ at $r = 64$) would be informative.

(L4) Quantization confound (addressed by C3). The original concern was that 4-bit NF4 quantization might suppress singular-value growth, making our sub-threshold SE observation an artefact of precision rather than language-specific dynamics. The C3 BF16 control (Kyrgyz, $r=16$, $lr=2 \times 10^{-4}$, 3 epochs, identical to E3 except for full-precision base model) refutes this: BF16 SE behaviour is statistically indistinguishable from 4-bit E3 at every comparable point—final max SE 0.32 vs. 0.33 ± 0.01 , final mean SE 0.10 in both, and in-training peak max SE 0.47 vs. 0.45 ± 0.02 (Table 11). Frobenius growth, KY PPL, NER F1, log-likelihood span typing, and TUMLU accuracy are likewise within seed-level variance of E3

(Section E). Quantization is therefore not artificially suppressing SE—and is not the explanation for the SE-threshold falsification. The residual gap is that we ran C3 only at $n=1$ (cloud GPU cost); a second BF16 seed would tighten the comparison further.

(L5) Corpus imbalance between Turkic languages (partially addressed by E1b). We match corpora by raw byte size (~ 150 MB each), yet token counts differ substantially after tokenization (KZ: 14.1M, UZ: 5.4M, KY: 4.4M). To isolate tokenizer-coverage effects from raw-data-volume effects, we ran a small-corpus Kazakh control E1b on ~ 1.5 M training tokens (the corpus was sized via a word $\times$ fragmentation-ratio estimate aiming for the KY 4.4M target; the actual chunked-token count came out lower). Result: target KZ PPL rises from 2.73 (E1) to 4.17 (E1b) with the smaller corpus, confirming that data volume materially contributes to the E1 advantage. Importantly, E1b’s 4.17 remains lower than the Kyrgyz baseline’s 4.78 even though E1b trained on $\sim 3\times$ fewer tokens than KY—so Kazakh also enjoys a genuine tokenizer-/pretraining-coverage advantage that is not attributable to corpus size. Both factors contribute; neither alone explains the gap. A parallel control on Uzbek data, and a true 4.4M-token Kazakh control, would tighten this further; we defer them to future work.

(L6) NER as a primary metric. Our NER evaluation uses 3-shot prompting with $n=100$ samples per language and F1 derived from parse-success output. This is closer to a test of instruction-following than of entity knowledge. The $r=64$ “NER F1 = 0” result, in particular, is driven by 100% parse failures; the model may retain some entity knowledge that this evaluation cannot access. Log-likelihood-based NER scoring (scoring candidate entity spans by model probability) would separate instruction loss from knowledge loss and is a priority follow-up.

(L7) Scope of generalization. All experiments use a single base model (Gemma-2-9B); results may differ for Llama 3, Mistral, or other architectures. We examine only three of the approximately 35 Turkic languages, excluding Turkish (the highest-resource Turkic language). Corpus size is fixed

at ~ 150 MB per language, precluding scaling analyses. SVD monitoring focuses on `lora_B`; `lora_A` dynamics may reveal complementary patterns we do not capture.

(L8) Post-hoc Frobenius observations. The growth ranges in Section 4.4 are derived from the same experiments on which their interpretation is built; they are observations requiring held-out validation, not established thresholds, and we deliberately reframe them as such in Section 4.4.

(L9) Uzbek script coverage. Our Uzbek training corpus is restricted to Cyrillic-script sources, while WikiANN-uz uses primarily Latin script: this is the most likely explanation for the 2% entity-string overlap reported in L2 (vs. $\sim 60\%$ for KZ/KY where script alignment is uniform Cyrillic). Since Uzbekistan officially uses Latin script and most contemporary formal text is Latin-based, our Uzbek perplexity and target-language NER results should not be generalized to Latin-script data without a separate evaluation. The script-mismatch nonetheless gives us an unintended benefit: Uzbek serves as a cleaner cross-lingual probe than KZ or KY because memorization-driven inflation is negligible, so the strong transfer-induced UZ improvements (e.g., E1’s 0.578 F1) are difficult to attribute to data leakage.

7 Conclusion

We have presented a systematic study of the spectral dynamics of LoRA adapters during fine-tuning on three low-resource Turkic languages. Our experiments yield three principal findings.

First, the spectral-energy threshold of [Biderman et al. \(2024\)](#) is **not predictive** in our regime. The E5/E5b/E5c three-way decomposition shows that at matched rank ($r=64$), three runs converging to near-identical settled spectral trajectories (final max SE 0.06–0.12, final mean SE ~ 0.03 in all three) span the full functional range from catastrophic collapse (E5: F1= 0, TUMLU 23%) to baseline-better target performance (E5c: KY PPL 3.86 vs. 4.78 for $r=16$). SE cannot detect the difference. The C3 BF16 control rules out 4-bit quantization as the cause: BF16 SE behaviour is statistically indistinguishable from

4-bit E3 (Table 11), so the sub-threshold values are not an artefact of precision. Our contribution here is not a new failure mode but an empirical demonstration that the SE diagnostic silently misses an ordinary learning-rate-induced collapse in the low-resource quantized regime where it is most likely to be deployed.

Second, **Frobenius norm growth** is a useful scale-aware monitoring signal but *not* a universal diagnostic. Within underrepresented languages it reliably separates regimes ($1.13\times$ for transfer vs. $15\text{--}38\times$ for high-LR $r=64$). Across language categories it fails: our English control absorbs $33\times$ norm growth with minimal forgetting, comparable to Kazakh’s $38\times$ *with* forgetting. We therefore recommend pairing $\|\mathbf{B}\|_F$ trajectory monitoring with cross-lingual PPL probes rather than relying on fixed thresholds.

Third, consistently across our two seeds, **warm-starting from a related-language checkpoint preserves source-language knowledge**: Kazakh→Kyrgyz transfer halves cross-lingual Kazakh PPL (23.5 ± 0.1 vs. 50.0 ± 2.4 under direct training—a $2\times$ effect that is large relative to the seed-to-seed spread observed in either condition) while requiring only +13% Frobenius-norm growth (vs. $8\times$ for direct training). The single-seed NER-F1 improvement (+61%) does *not* reproduce at seed 123 (−2%); we therefore characterize the transfer effect primarily as *retention* of the source language rather than target-language improvement (Section F). A cross-family control (E6b: EN→KY, $n=1$) rules out a “generic warm-start” explanation: starting from English instead of Kazakh yields KZ PPL 47.46, essentially identical to direct training (47.58) and far above the related-language transfer’s 23.49. Related-language warm-starts are therefore qualitatively different from unrelated warm-starts in this regime; a second-seed E6b replication remains a priority follow-up.

These findings have immediate practical implications for the growing community of researchers and practitioners adapting LLMs to low-resource languages. Future work should validate our findings across diverse model architectures (Llama-3, Mistral, etc.; our BF16 control C3 already rules out 4-bit quantization as a confound), extend the analysis to ad-

ditional Turkic languages including Turkish, compare alternative LoRA initializations such as PiSSA (Meng et al., 2024) that may exhibit different spectral dynamics, investigate directional norm regularization as a training objective, and explore adaptive rank selection strategies informed by real-time spectral monitoring. An open question is whether `lora_A` matrices exhibit complementary spectral signals when `lora_B` norms explode—joint analysis of both matrices may yield a more complete diagnostic.

Ethics Statement

This work involves fine-tuning language models on text corpora from three Turkic languages and English. While we aim to improve NLP capabilities for underserved language communities, we acknowledge that language models can amplify biases present in training data. Our training corpora, drawn from Wikipedia, literary texts, and web sources, may contain cultural biases or factual inaccuracies.

Data availability: The Kazakh and Uzbek corpora are compiled from publicly available sources (Wikipedia, open book collections). The Kyrgyz corpus is curated from licensed literary and encyclopedic sources and cannot be redistributed due to copyright restrictions. In the spirit of open science, we release all experimental code, SVD monitoring logs, and evaluation results via an anonymous repository.³ Trained LoRA adapters will be publicly released upon acceptance to preserve author anonymity during review.

Acknowledgments

We thank the creators of the Kazakh, Kyrgyz, and Uzbek corpora used in this study, as well as the developers of the TUMLU benchmark for providing standardized Turkic language evaluation resources. All experiments were conducted on personal hardware without external funding.

Reproducibility Statement

All experimental code (training with SVD monitoring, evaluation, and plot-

³<https://anonymous.4open.science/r/Spectral-Collapse-Turkic-Conference-14D8>

ting scripts), SVD monitoring logs, and evaluation reports are publicly available at <https://anonymous.4open.science/r/Spectral-Collapse-Turkic-Conference-14D8>. Fine-tuned LoRA adapters for all eight experiments will be released on a public model hub upon acceptance to preserve author anonymity during review. The Kazakh, Uzbek, and English training corpora are derived from publicly available sources and can be reconstructed using the provided data preparation scripts. The Kyrgyz corpus cannot be redistributed due to licensing restrictions on its source materials; consequently, the KY training stage is not fully reproducible. However, the released KY adapters and SVD logs enable full reproduction of all evaluation and analysis stages.

References

- Armen Aghajanyan, Sonal Gupta, and Luke Zettlemoyer. 2021. Intrinsic dimensionality explains the effectiveness of language model fine-tuning. In *Proceedings of ACL-IJCNLP*.
- Orevaoghene Ahia, Julia Kreutzer, and Sara Hooker. 2023. Do all languages cost the same? tokenization in the era of commercial language models. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*.
- Dan Biderman, Jacob Portes, Jose Javier Gonzalez Ortiz, Mansheej Paul, Philip Greengard, Connor Jennings, Daniel King, Sam Havens, Vitaliy Chiley, Jonathan Frankle, Cody Blakeney, and John P Cunningham. 2024. LoRA learns less and forgets less. *Transactions on Machine Learning Research*.
- Zoltan Csaki, Bo Peng, Douglas Steiner, and 1 others. 2024. Sambalingo: Teaching large language models new languages. In *arXiv preprint arXiv:2404.05829*.
- Tim Dettmers, Artidoro Pagnoni, Ari Holtzman, and Luke Zettlemoyer. 2023. QLoRA: Efficient finetuning of quantized language models. In *Advances in Neural Information Processing Systems*.
- Gemma Team, Morgane Riviere, Shreya Pathak, Pier Giuseppe Sessa, and 1 others. 2024. Gemma 2: Improving open language models at a practical size. *arXiv preprint arXiv:2408.00118*.
- Edward J Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yuanzhi Li, Shean Wang, Lu Wang, and Weizhu Chen. 2022. LoRA: Low-rank adaptation of large language models. In *International Conference on Learning Representations*.
- Maral Isbarov and 1 others. 2024. TUMLU: Turkic multi-task language understanding benchmark. In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*.
- Uijeong Jang, Jason D Lee, and Ernest K Ryu. 2024. Lora training in the NTK regime has no spurious local minima. In *International Conference on Machine Learning*.
- James Kirkpatrick, Razvan Pascanu, Neil Rabinowitz, Oriol Vinyals, Guillaume Desjardins, Andrei A Rusu, Kieran Milan, John Quan, Tiago Ramalho, Agnieszka Grabska-Barwinska, and 1 others. 2017. Overcoming catastrophic forgetting in neural networks. In *Proceedings of the National Academy of Sciences*, volume 114, pages 3521–3526.
- Suhas Kotha, Jacob Mitchell Springer, and Aditi Raghunathan. 2024. Understanding catastrophic forgetting in language models via implicit inference. In *Proceedings of ICLR*.
- Vladislav Lialin, Namrata Deshpande, and Anna Rumshisky. 2023. ReLoRA: High-rank training through low-rank updates. In *NeurIPS 2023 Workshop on Efficient Natural Language and Speech Processing*.
- Shih-Yang Liu, Chien-Yi Wang, Hongxu Yin, Pavlo Molchanov, Yu-Chiang Frank Wang, Kwang-Ting Cheng, and Min-Hung Chen. 2024. DoRA: Weight-decomposed low-rank adaptation. In *Proceedings of ICML*.
- Yun Luo, Zhicheng Yang, Fandong Meng, Yafu Li, Jie Zhou, and Yue Zhang. 2023. An empirical study of catastrophic forgetting in large language models during continual fine-tuning. In *arXiv preprint arXiv:2308.08747*.
- Sadhika Malladi, Tianyu Gao, Eshaan Nichani, Alex Damian, Jason D Lee, Danqi Chen, and Sanjeev Arora. 2023. A kernel-based view of language model fine-tuning. In *International Conference on Machine Learning*.
- Charles H Martin and Michael W Mahoney. 2021. Implicit self-regularization in deep neural networks: Evidence from random matrix theory and implications for training. In *Journal of Machine Learning Research*, volume 22, pages 1–73.
- Michael McCloskey and Neal J Cohen. 1989. Catastrophic interference in connectionist networks: The sequential learning problem. *Psychology of Learning and Motivation*, 24:109–165.
- Fanxu Meng, Zhaohui Wang, and Muhan Zhang. 2024. PiSSA: Principal singular values and singular vectors adaptation of large language models. *Advances in Neural Information Processing Systems*.

Xiaoman Pan, Boliang Zhang, Jonathan May, Joel Nothman, Kevin Knight, and Heng Ji. 2017. Cross-lingual name tagging and linking for 282 languages. In *Proceedings of the 55th Annual Meeting of the Association for Computational Linguistics*.

Jonas Pfeiffer, Ivan Vulić, Iryna Gurevych, and Sebastian Ruder. 2020. MAD-X: An adapter-based framework for multi-task cross-lingual transfer. In *Proceedings of EMNLP*.

Olivier Roy and Martin Vetterli. 2007. The effective rank: A measure of effective dimensionality. *European Signal Processing Conference*, pages 606–610.

Phillip Rust, Jonas Pfeiffer, Ivan Vulić, Sebastian Ruder, and Iryna Gurevych. 2021. How good is your tokenizer? on the monolingual performance of multilingual language models. In *Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics*.

Ruslan Yeshpanov, Yerbolat Khassanov, and Huseyin Atakan Varol. 2024. Kazllm: A study of large language models for kazakh language processing. *arXiv preprint arXiv:2407.15454*.

Rustem Yeshpanov, Yerbolat Khassanov, and Huseyin Atakan Varol. 2022. KazNERD: Kazakh named entity recognition dataset. *arXiv preprint arXiv:2111.13419*.

Qingru Zhang, Minshuo Chen, Alexander Bukharin, Pengcheng He, Yu Cheng, Weizhu Chen, and Tuo Zhao. 2023. AdaLoRA: Adaptive budget allocation for parameter-efficient fine-tuning. In *International Conference on Learning Representations*.

A Detailed Training Configurations

All experiments were conducted on a single NVIDIA GeForce RTX 5080 (16 GB) using PyTorch 2.10, Transformers 5.2, and PEFT 0.18. Table 7 provides the full training statistics.

B Effective Rank Dynamics

Figure 8 shows the effective rank dynamics across experiments. The $r=64$ configuration shows increasing effective rank (from ~ 6 to ~ 26), indicating growing parameter utilization. The diverged experiment (E7) is the only configuration where effective rank *decreases* (from 43.3 to 37.1), consistent with true spectral collapse.

E7 divergence analysis. The diverged experiment ($r=64$, $\text{lr}=10^{-3}$) provides the most extreme datapoint. Training terminated at

step 800 due to loss divergence. Even at this point, maximum SE was only 0.18—still far below the 0.7 threshold—while mean Frobenius norm had already reached 6.75 ($19\times$ growth in just 800 steps, compared to $15.6\times$ over 4,885 steps for E5). The effective rank decrease ($43.3 \rightarrow 37.1$) in the final 200 steps is the only instance of true rank contraction in our experiments, suggesting that SE-based collapse detection may only trigger *after* the model has already diverged—too late for practical intervention. The Frobenius norm, by contrast, shows clear exponential growth from step 100 onward, providing a much earlier warning signal.

C Log-Likelihood NER: Type Classification under Gold Spans

To separate loss of entity knowledge from loss of instruction-following/output-formatting, we add a second NER evaluation that does not require the model to generate structured output. Given the gold-span boundaries from WikiANN, we present each span in context and score three candidate completions (“person”, “organization”, “location”) by log-likelihood under the adapter; the predicted type is the argmax. Because this evaluation never relies on the model’s output format, it is immune to the 100% parse-failure mode that drives E5’s NER $F1 = 0$ under generation-based scoring.

Table 8 reports per-language type accuracy and macro-F1 on the same 100-example test sets used for Table 3.

D WikiANN–Training Entity Overlap

We quantify the entity-level overlap between each language’s training corpus and the WikiANN test split it is evaluated on. For every gold span in the WikiANN test set we test whether the entity string appears verbatim (case-insensitive) anywhere in the corresponding training corpus. The unique-overlap rate is an upper bound on the fraction of NER F1 that could be explained by memorization of seen entity surface forms.

ID	Experiment	Train Loss	Eval Loss	Eval PPL	Tokens	Time (min)
E1	KZ Baseline (14.1M)	0.275	1.259	3.52	14.1M	477
E1b	KZ Small-corpus (1.5M)	1.186	1.461	4.31	1.5M	122
E2	UZ Baseline	1.193	1.822	6.18	5.4M	341
E3	KY Baseline	0.815	1.920	6.82	4.4M	183
E4	KY Overfit	1.050	1.963	7.12	4.4M	1193
E5	KY $r=64$ high-LR	0.274	2.164	8.71	4.4M	157
E5b	KY $r=64$ std-LR (5ep)	1.278	1.964	7.13	4.4M	605
E5c	KY $r=64$ epoch-matched	1.652	1.898	6.67	4.4M	360
E6	KZ→KY	1.809	1.932	6.90	4.4M	355
E8	EN Control	1.821	2.017	7.51	12.7M	1351

Table 7: Full training statistics for all experiments. Eval loss and PPL are computed on the validation split during training (distinct from the post-hoc evaluation in Table 3). The higher train loss for E6 (KZ→KY) reflects the adapter starting from KZ-optimized weights that initially perform poorly on KY data.




Figure 8: Effective rank dynamics during training. The $r=64$ configuration shows increasing effective rank despite functional degradation, further evidence that standard spectral metrics can be misleading.

E BF16 vs. 4-bit Control (C3)

Table 11 reports a side-by-side comparison of the 4-bit E3 baseline (averaged across our two seeds where available) and the full-precision BF16 control (C3, single seed, run on a cloud A100 40GB instance). All hyperparameters except quantization are identical.

F Seed-Variance Replication

To provide at least partial seed-variance estimates for our most consequential configurations, we rerun two experiments under a second seed (123, while the originals used seed 42). Training hyperparameters and data are identical across seeds; only the random seed for initialization, data shuffling, and train/val

split differs.

These results sharpen the interpretive stance taken throughout the paper: we treat cross-lingual PPL differences as the most reliable quantitative signal, and the various NER scores as directional indicators whose per-experiment magnitudes should not be over-interpreted. A full $n=3$ replication of all eight experiments remains the most important reproducibility follow-up (see L1); the present two-seed data covers the two configurations whose comparison underlies our main transfer claim.

G Per-Type NER Results

Table 13 breaks down NER performance by entity type for the Kyrgyz evaluation.

Experiment	KY	KZ	UZLanguage	Total spans	Unique	Unique overlap
E1 KZ Baseline (14.1M tok)	57.7%	72.3%	70.9%Kazakh (KZ)	1,115	757	65.7% (497/757)
E1b KZ Small-corpus (1.5M)	46.0%	70.5%	69.9%Kyrgyz (KY)	111	93	55.9% (52/93)
E2 UZ Baseline	54.9%	71.4%	65.0%Uzbek (UZ)	1,011	602	2.2% (13/602)
E3 KY Baseline	50.4%	69.6%	63.1%			
E4 KY Overfit (10ep)	56.8%	71.4%	60.2%			
E5 KY $r=64$ (lr 5×10^{-4})	59.5%	50.0%	55.3%			
E5b KY $r=64$ 5ep (lr 2×10^{-4})	53.1%	70.5%	69.9%			
E5c KY $r=64$ 3ep (lr 2×10^{-4})	54.9%	67.9%	63.1%			
E6 KZ→KY	59.5%	73.2%	66.0%			
E8 EN Control	56.8%	73.2%	68.0%			

Table 8: Log-likelihood span-type accuracy across experiments (gold-span boundaries, argmax over {PER, ORG, LOC}). Random-baseline accuracy is 33.3%. The parse-failure-immune evaluation reveals that **E5 retains and even marginally exceeds target-language (KY) entity knowledge** (59.5%, tied for best with E6) despite scoring $F1 = 0$ under generation-based NER—confirming that the high-LR $r=64$ collapse destroys instruction-following, not entity knowledge itself. Cross-lingual degradation, however, does appear in E5 on non-target languages (KZ 50.0% vs. 70–73% for others), consistent with the cross-lingual PPL increase. Under this cleaner metric, the KZ→KY transfer advantage over the KY baseline is modest (59.5% vs. 50.4% on KY; +9.1 absolute percentage points) rather than the +61% relative improvement suggested by generation F1, reinforcing our framing of cross-lingual transfer as a *directional* rather than precise effect.

H TUMLU Accuracy by Language

Table 14 presents TUMLU accuracy scores across all languages and experiments.

I Functional Collapse: Raw Output Examples

To illustrate the nature of functional collapse at $r=64$, we present raw model outputs for two tasks. The model completely ignores instructions and exhibits degenerate repetition.

NER prompt: Extract named entities from the text. Format each as: ENTITY ||| TYPE

Input: Бишкекте Кыргыз Республикасынын Президентин Садыр Жапаров менен Казакстандын Президентин Касым-Жомарт Токаев жолукту.

$r=64$ **output:** Кыргыз Республикасынын Президентинин кызматына кандидат болгон. Кыргыз Республикасынын Президентинин

Table 9: Verbatim case-insensitive overlap between WikiANN test entities and our training corpora. KZ and KY show high overlap (memorization upper bound $\sim 60\%$); UZ shows almost none, which we attribute to script mismatch between our Cyrillic UZ corpus and Latin-script WikiANN-uz (L9). The UZ low overlap incidentally turns Uzbek into a cleaner cross-lingual transfer probe: improvements in UZ NER are unlikely to be memorization artefacts.

Language	PER	ORG	LOC
Kazakh	71% (221/313)	62% (170/276)	63% (106/168)
Kyrgyz	33% (11/33)	57% (17/30)	80% (24/30)
Uzbek	3% (2/67)	4% (8/202)	1% (3/333)

Table 10: Per-type unique overlap. The KY pattern (LOC $\gg$ PER) is consistent with location names being a comparatively closed-class compared to person names. UZ remains uniformly low across types, confirming the script-mismatch explanation.

кызматына кандидат болгон. Кыргыз Республикасынын Президентинин кызматына кандидат болгон... (repeats indefinitely)

Text continuation **prompt:** Кыргызстан --- это

$r=64$ **output:** не только, не только, не только, не только... (single bigram repeated 25+ times until max tokens)

Both examples demonstrate that the model has lost instruction-following capability entirely. Rather than extracting entities or generating coherent text, it produces degenerate repetitions—a hallmark of functional collapse distinct from spectral collapse.

Metric	E3 (4-bit, 2 seeds)	C3 (BF16, $n=1$)
KY PPL	4.81 ± 0.02	4.99
KZ PPL (cross-ling)	50.0 ± 2.4	43.97
UZ PPL (cross-ling)	55.5 ± 1.4	59.02
NER F1 KY (gen)	0.170 ± 0.013	0.164
NER TypeAcc KY	55.0 ± 4.6	58.6
TUMLU KY	36.1 ± 1.4	36.1
$\ B\ _F$ growth	$8.36\times$	$8.60\times$
Final SE max	0.33 ± 0.01	0.32
Final SE mean	0.10	0.10
Peak SE max (training)	0.45 ± 0.02	0.47

Table 11: 4-bit NF4 vs. BF16 control on the same KY baseline configuration (rank 16, lr 2×10^{-4} , 3 epochs). All downstream metrics fall within or near the seed-level variance observed for the 4-bit run. Frobenius norm growth is essentially identical. Spectral energy in BF16 is statistically indistinguishable from 4-bit at every comparable point (final max, final mean, in-training peak), which still rules out the original concern: 4-bit quantization is not artificially suppressing SE. The SE-threshold falsification therefore carries over to full precision. We did not run a second seed for C3 due to compute cost (the BF16 run was completed on rented cloud GPU); $n=1$ is acknowledged in L1.

Metric	E1b		E3		E5		E5b		E5c		E6		E8 (EN ctrl)	
	s=42	s=123	s=42	s=123	s=42	s=123	s=42	s=123	s=42	s=123	s=42	s=123	s=42	s=123
KY PPL ↓	25.55	28.81	4.78	4.83	5.90	6.01	4.55	4.62	3.86	3.92	4.73	4.78	14.96	14.82
KZ PPL ↓	4.17	4.15	47.58	52.47	659.25	606.81	128.04	107.23	93.52	98.62	23.49	23.61	5.54	5.50
UZ PPL ↓	23.17	21.93	56.85	54.05	742.73	763.79	162.68	139.89	111.25	96.61	124.19	114.54	16.55	16.28
NER F1 KY (gen) ↑	0.150	0.139	0.157	0.182	0.000	0.000	0.146	0.116	0.115	0.169	0.253	0.179	0.209	0.227
TypeAcc KY ↑	46.0	53.1	50.4	59.5	59.5	53.1	53.1	64.9	54.9	64.9	59.5	64.9	56.8	54.9
TypeAcc KZ ↑	70.5	70.5	69.6	70.5	50.0	40.2	70.5	71.4	67.9	69.6	73.2	75.0	73.2	72.3
TUMLU KY ↑	37.3	38.6	34.7	37.4	23.2	23.9	33.2	33.2	38.5	38.0	33.7	34.9	37.7	36.2
$\ B\ _F$ growth	—	2.9×	8.4×	—	15.6×	15.5×	18.2×	—	8.7×	8.7×	1.1×	—	33.4×	33.2×

Table 12: Two-seed replication for the seven configurations underlying our main claims: small-corpus KZ baseline (E1b, 1.5M tokens), KY baseline (E3), catastrophic collapse (E5), rank-LR control (E5b), epoch-matched control (E5c), KZ→KY transfer (E6), and EN control (E8). All TypeAcc values are percentages. **Five robust claims emerge across seeds:** (i) *KZ-retention via transfer halves PPL* (23.5 ± 0.1 vs. 50.0 ± 2.4 for direct training, $10\times$ outside seed variance); (ii) *$r=64$ outperforms $r=16$ baseline at matched conditions* (E5c KY PPL 3.89 ± 0.03 vs. E3 4.81 ± 0.03); (iii) *rank-LR disentanglement is robust* (E5b sustains $F1 \geq 0.12$, TUMLU $\sim 33\%$ in both seeds, far outside E5’s $F1 = 0$, TUMLU $\sim 23\%$); (iv) *E1b data-vs-coverage finding is robust* (E1b’s KZ PPL 4.16 ± 0.01 across seeds is reliably below E3’s 4.81 despite $\sim 3\times$ fewer training tokens, while well above E1’s 2.73 at full corpus); (v) *the EN-vs-Turkic Frobenius dissociation is robust* (E8 $33.3 \pm 0.1\times$ growth with KZ PPL 5.52 ± 0.02 , indistinguishable from E1’s $38.5\times$ growth with KZ PPL ~ 48 across seeds: same norm magnitude, opposite forgetting outcomes—supporting our claim that the Frobenius diagnostic fails across language categories). The single-seed transfer NER F1 headline ($+61\%$) does *not* survive reseeding (-2% at seed 123); we therefore reframe transfer as a retention mechanism.

Experiment	PER F1	ORG F1	LOC F1
KZ Baseline 14.1M (on KY)	0.236	0.162	0.229
KZ Small-corpus 1.5M (on KY)	0.101	0.207	0.261
UZ Baseline (on KY)	0.095	0.150	0.209
KY Baseline	0.101	0.160	0.266
KY Overfit	0.224	0.255	0.119
KY $r=64$ (E5, lr 5×10^{-4})	0.000	0.000	0.000
KY $r=64$ (E5b, 5ep)	0.126	0.178	0.177
KY $r=64$ (E5c, 3ep)	0.107	0.139	0.114
KZ→KY	0.344	0.202	0.214

Table 13: Per-type NER F1 scores on Kyrgyz WikiANN. The KZ→KY transfer achieves the highest PER F1 (0.344), substantially outperforming the KY baseline (0.101). The $r=64$ configuration produces zero scores across all types due to 100% parse failures.

Experiment	KY	KZ	UZ
KZ Baseline (14.1M tok)	35.7%	38.4%	32.8%
KZ Small-corpus (1.5M tok)	37.3%	35.6%	32.8%
UZ Baseline	31.9%	33.1%	35.4%
KY Baseline	34.7%	30.2%	29.6%
KY Overfit	33.9%	27.8%	31.5%
KY $r=64$ (E5, lr 5×10^{-4})	23.2%	22.2%	22.9%
KY $r=64$ (E5b, 5ep)	33.2%	27.7%	26.6%
KY $r=64$ (E5c, 3ep)	38.5%	29.8%	31.9%
KZ→KY	33.7%	31.7%	30.8%

Table 14: TUMLU accuracy across all experiments and languages. The KZ baseline achieves the highest accuracy on its target language (38.4%). The $r=64$ configuration drops to near-random levels ($\sim 23\%$ on 4-option multiple choice) across all languages.